





Support Vector Machine (SVM)




The underlying philosophy of SVM classifiers



- Find lines (hyperplanes) that correctly classify the training data
- Among all such lines (hyperplanes), pick the one that has the greatest distance to the points closest to it.

Margin (gap between decision boundary and hyperplanes)



“Decision Boundary”


A linear classifier’s decision boundary is also called a “hyperplane”. Hyperplane is a $(n-1)$ -dimensional subspace, for example, a line for observations in 2D space and a plane for observations in a 3D space, separating different classes of data observed in a n -dimensional space.

“Margin”


Margin is defined as the perpendicular distance of the nearest data points to the hyperplane. There are “hard margin” and “soft margin”. Our objective is to find the hyperplane with the maximum margin.

“Support Vector”


The nearest data points on the edge of margins are support vectors.



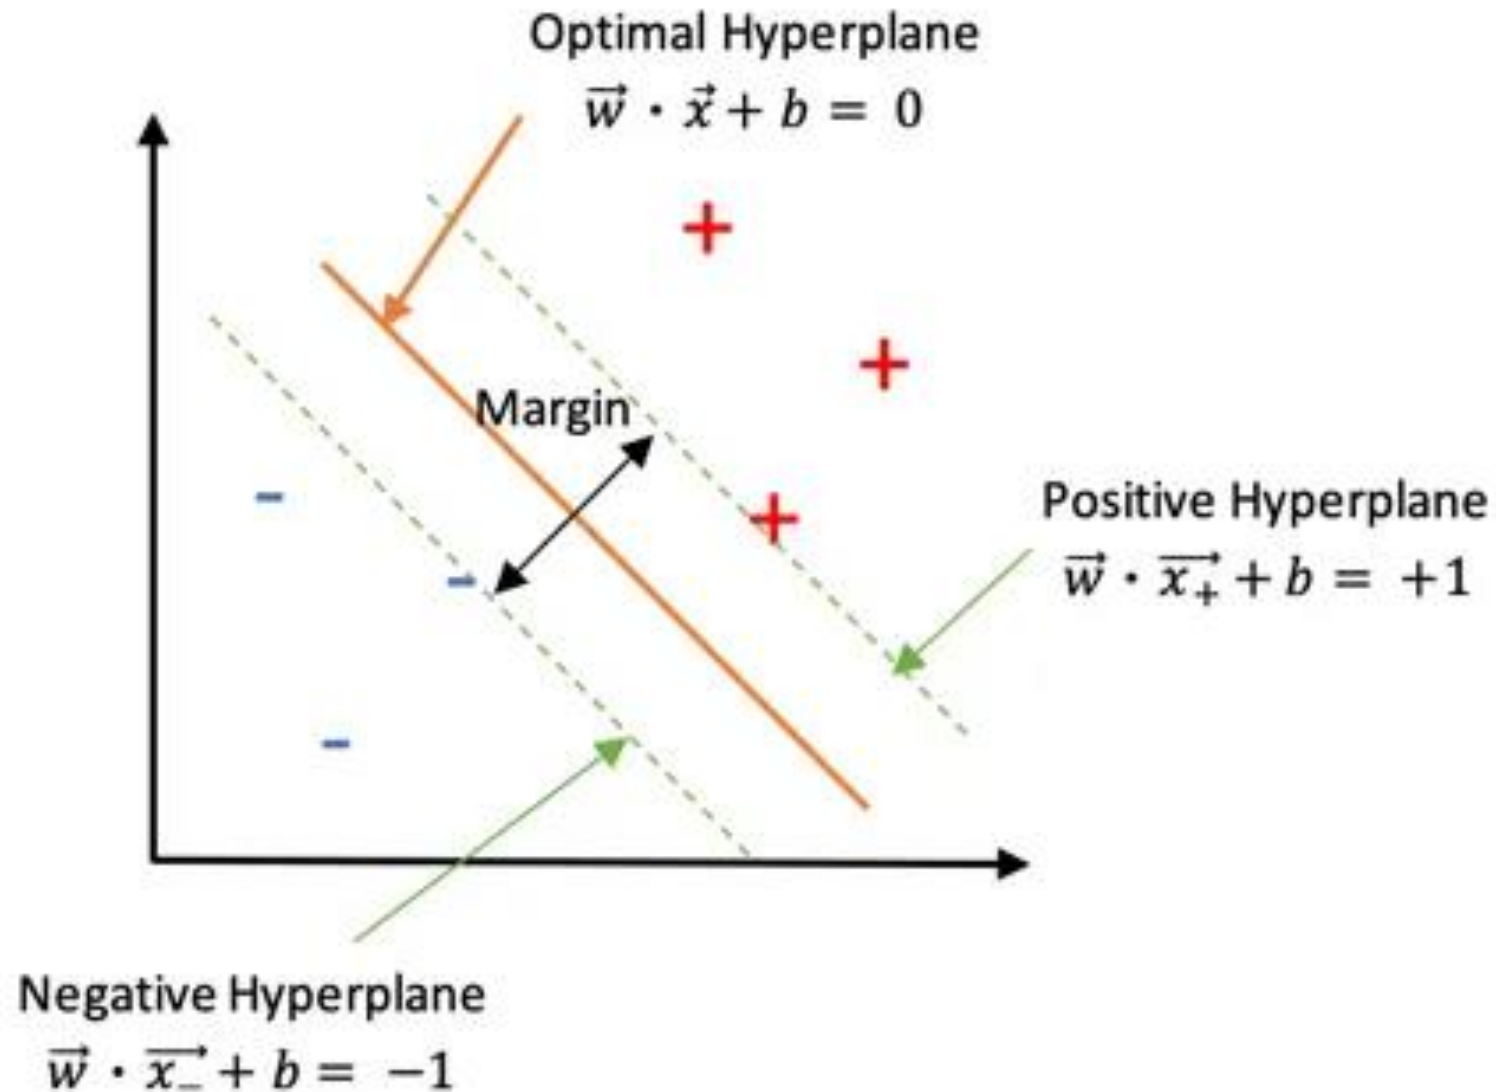
A hyperplane in \mathbb{R}^2 is a line



A hyperplane in \mathbb{R}^3 is a plane




Optimal hard margin classifier



Geometric Intuition and basic understanding

- W (parameters of the model) is a vector (geometric) perpendicular to hyperplane $\mathbf{w}^T \mathbf{X} + \mathbf{b} = 0$
- b is the offset of the hyperplane in vector space
- Our objective to train SVM is to find w and b which maximize the margin.



$$\frac{w}{\|w\|} \cdot (x_2 - x_1) = \text{width} = \frac{2}{\|w\|}$$


$$w \cdot x_2 + b = 1$$

$$w \cdot x_1 + b = -1$$

$$w \cdot x_2 + b - w \cdot x_1 - b = 1 - (-1)$$

$$w \cdot x_2 - w \cdot x_1 = 2$$

$$\frac{w}{\|w\|} (x_2 - x_1) = \frac{2}{\|w\|}$$



$$\max \frac{2}{\|\vec{w}\|}$$

s.t.

$(\vec{w} \cdot \vec{x} + b) \geq 1, \forall \vec{x}$ of class 1

$(\vec{w} \cdot \vec{x} + b) \leq -1, \forall \vec{x}$ of class 2


Mathematical intuition of Support Vector Machine


We find w and b by solving the following objective function using Quadratic Programming.

Constrained optimization problem

$$\begin{aligned} \min_{w, b} \quad & \frac{1}{2} w^T w \\ \text{subject to} \quad & y_i \cdot (w^T x_i + b) \geq 1, i = 1, \dots, n \end{aligned}$$

When a hard margin classifier is not applicable





slack variable:

$$\xi_i$$

Allow some instances to fall off the margin, but penalize them

correctly classified points
must not be penalized

misclassified points inside the
margin, or on the wrong side of the
margin must be penalized

correctly classified points
must not be penalized

Measure the misclassification error for each training example

$$\xi_i = \begin{cases} 0, & y_i \cdot (\mathbf{w}^T \mathbf{x}_i + b) \geq 1 \\ y_i \cdot (\mathbf{w}^T \mathbf{x}_i + b) & y_i \cdot (\mathbf{w}^T \mathbf{x}_i + b) < 1 \end{cases}$$

Penalize each **misclassification by the size of the violation**, using the **hinge loss** (contrast with the loss function of the Perceptron)


$$\xi_i = L(f(\mathbf{x}_i), y_i) = \max\{0, 1 - y_i \cdot (\mathbf{w}^T \mathbf{x}_i + b)\}$$

Measure the misclassification error for each training example

$$\xi_i = \begin{cases} 0, & y_i \cdot (\mathbf{w}^T \mathbf{x}_i + b) \geq 1 \\ y_i \cdot (\mathbf{w}^T \mathbf{x}_i + b) & y_i \cdot (\mathbf{w}^T \mathbf{x}_i + b) < 1 \end{cases}$$

Penalize each misclassification by the size of the violation, using the hinge loss (contrast with the loss function of the Perceptron)

$$\xi_i = L(f(\mathbf{x}_i), y_i) = \max\{0, 1 - y_i \cdot (\mathbf{w}^T \mathbf{x}_i + b)\}$$



Constraint becomes :

$$y_i(w \cdot x_i + b) \geq 1 - \xi_i, \quad \forall x_i$$





$$\xi_i \geq 0$$

Objective function


penalizes for misclassified instances and those within the margin

$$\min \frac{1}{2} \|w\|^2 + C \sum_i \xi_i$$

C trades-off margin width and misclassifications



Support Vector Machine Classifier




- Hyperplane
- Support Vectors
- Margin
- Linearly separable data

Soft-Margin SVM: Primal Problem

Problem: Find a linear classifier $f(x) = \mathbf{w}^T \mathbf{x} + b$ with the largest margin such that

$\text{sign}(f(x)) = +1$, when positive example

$\text{sign}(f(x)) = -1$, when negative example
and misclassifications are minimized.



soft-margin support vector machine


$$\begin{aligned} \min \quad & \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \xi_i \\ \text{s.t.} \quad & y_i (\mathbf{w}' \mathbf{x}_i - b) \geq 1 - \xi_i \quad \forall i = 1 \dots n \\ & \xi_i \geq 0 \end{aligned}$$

This model is called the **soft-margin SVM** as it softens the classification constraints with **slack variables** (ξ_i) and allows flexibility for misclassifications by the model

hard-margin support vector machine

$$\begin{aligned} \min \quad & \frac{1}{2} \|\mathbf{w}\|_2^2 \\ \text{s.t.} \quad & y_i (\mathbf{w}' \mathbf{x}_i - b) \geq 1 \quad \forall i = 1 \dots n \end{aligned}$$

Soft-Margin SVM: Dual Problem



soft-margin svm dual

$$\begin{aligned} \max \quad & -\frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \mathbf{x}'_i \mathbf{x}_j + \sum_{i=1}^n \alpha_i \\ \text{s.t.} \quad & \sum_{i=1}^n \alpha_i y_i = 0 \\ & 0 \leq \alpha_i \leq C, \quad \forall i = 1 \dots n \end{aligned}$$


the only difference between the soft-margin and hard-margin SVM dual problems is that the Lagrange multipliers (α_i training example weights) are upper-bounded by the **regularization parameter** ($0 \leq \alpha_i \leq C$)

hard-margin svm dual



$$\begin{aligned} \max \quad & -\frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \mathbf{x}'_i \mathbf{x}_j + \sum_{i=1}^n \alpha_i \\ \text{s.t.} \quad & \sum_{i=1}^n \alpha_i y_i = 0, \\ & \alpha_i \geq 0, \quad \forall i = 1 \dots n \end{aligned}$$

$0 \leq \alpha_i \leq \infty$

Kernel Trick in Support Vector Classification



Motivation For Kernels




SVM Kernels

Feature (x)	-6	-5	-4	-3	-2	-1	0	1	2	3	4	5	6
-------------	----	----	----	----	----	----	---	---	---	---	---	---	---

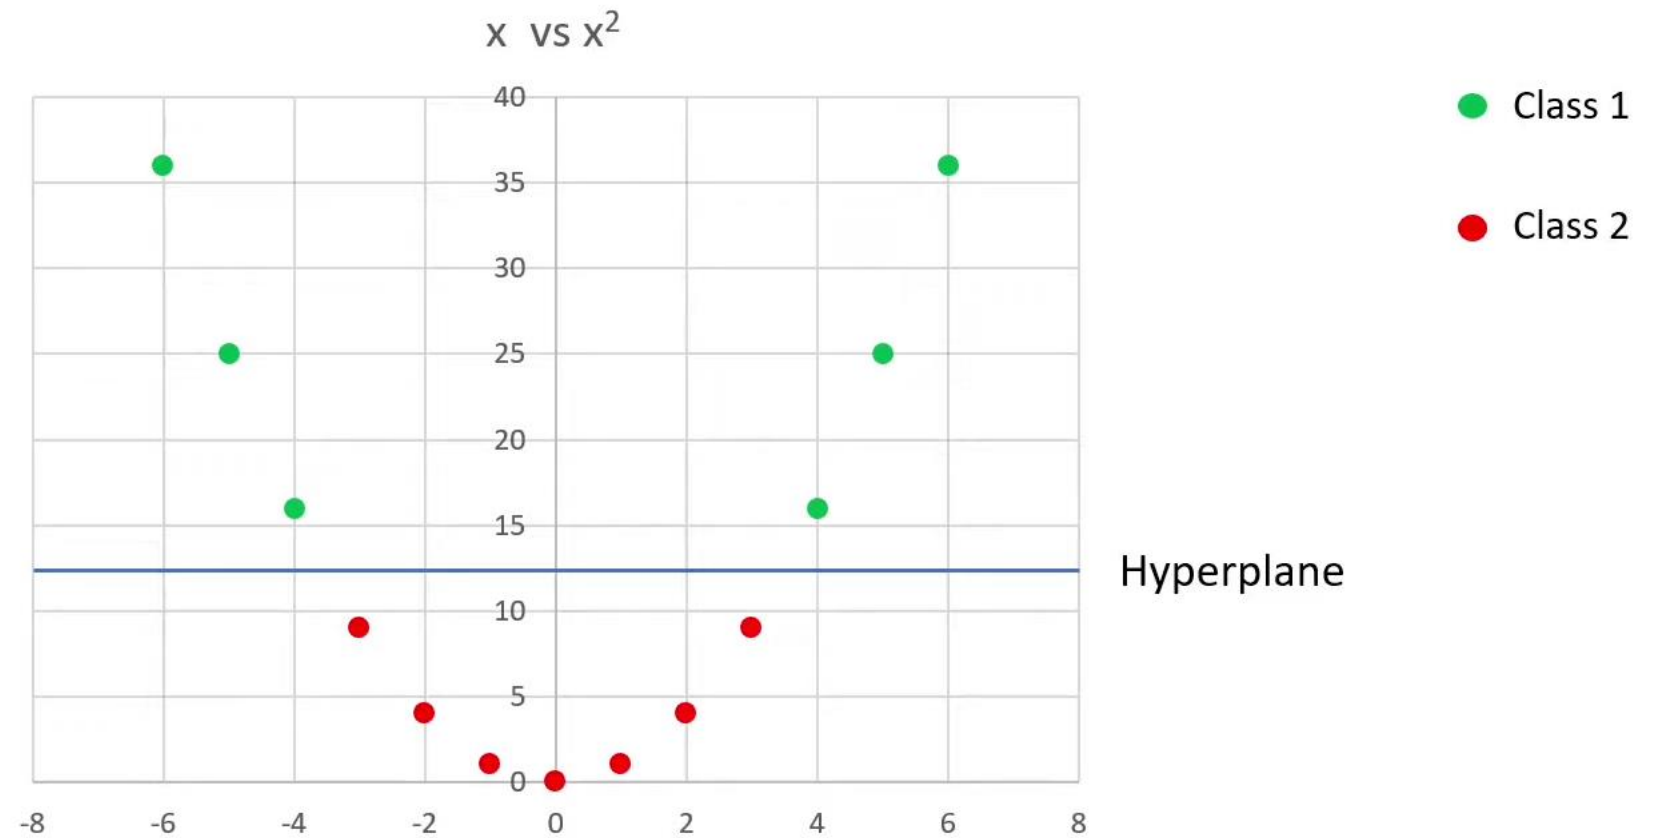
● Class 1

● Class 2





SVM Kernels

Feature (x)	-6	-5	-4	-3	-2	-1	0	1	2	3	4	5	6
x^2	36	25	16	9	4	1	0	1	4	9	16	25	36







Guess the feature transformation function



Guess the feature transformation function

Here we apply the transformation $\phi(x) = x \bmod 2$





$$\phi : \mathbb{R}^2 \rightarrow \mathbb{R}^3$$

$\phi : \mathcal{X} \rightarrow F$ *Second-degree polynomial mapping*

$$\phi(\mathbf{x}) = \phi\left(\begin{pmatrix} x_1 \\ x_2 \end{pmatrix}\right) = \begin{pmatrix} x_1^2 \\ \sqrt{2} x_1 x_2 \\ x_2^2 \end{pmatrix}$$


Example

$$x = \begin{bmatrix} x_1 & x_2 \\ 1 & 1 \\ 1 & 2 \\ 1 & 3 \\ 2 & 1 \end{bmatrix} \quad \phi(x) = \begin{bmatrix} x_1 & x_2 & x_3 \\ 1 & 1 & \sqrt{2} \\ 1 & 4 & \sqrt{2} \cdot 2 \\ 1 & 9 & \sqrt{2} \cdot 3 \\ 4 & 1 & 2 \end{bmatrix}$$

Take away points.....

- We have seen how higher dimensional transformations can allow us to separate data in order to make classification predictions. It seems that in order to train a support vector classifier and optimize our objective function, we would have to perform operations with the higher dimensional vectors in the transformed feature space.
- In real applications, there might be many features in the data and applying transformations that involve many polynomial combinations of these features will lead to extremely high and impractical computational costs.

Non-linear SVMs: Feature Space




What are the steps for nonlinear SVM classification?

- Feature transformation in higher dimensional space
- Train linear SVM classifier and find out the hyperplane
- During test time also use the transformation function on test input features and do the inference.

Nonlinear SVM Classifiers

Solution Approach 1 (Explicit Transformation)

- transform data to a higher dimensional feature space
- train linear SVM classifier in the high-dim. space
- transform high-dimensional linear classifier back to the original space to obtain a nonlinear classifier




if we have m training features, the size of the transformation grows very fast; **explicit transformations** can become **very expensive**

$$\Phi(\mathbf{x}) = \begin{pmatrix} 1 \\ \sqrt{2}x_1 \\ \sqrt{2}x_2 \\ \vdots \\ \sqrt{2}x_m \\ x_1^2 \\ x_2^2 \\ \vdots \\ x_m^2 \\ \sqrt{2}x_1x_2 \\ \sqrt{2}x_1x_3 \\ \vdots \\ \sqrt{2}x_1x_m \\ \sqrt{2}x_2x_3 \\ \vdots \\ \sqrt{2}x_1x_m \\ \vdots \\ \sqrt{2}x_{m-1}x_m \end{pmatrix}$$

Annotations for the terms in $\Phi(\mathbf{x})$:

- Constant Term (1)
- Linear Terms ($\sqrt{2}x_1, \sqrt{2}x_2, \dots, \sqrt{2}x_m$)
- Pure Quadratic Terms ($x_1^2, x_2^2, \dots, x_m^2$)
- Quadratic Cross-Terms ($\sqrt{2}x_1x_2, \sqrt{2}x_1x_3, \dots, \sqrt{2}x_1x_m, \sqrt{2}x_2x_3, \dots, \sqrt{2}x_{m-1}x_m$)

The Kernel Trick




data in higher-dimensional space

$$\phi(\mathbf{x}) = (x_1^2, \sqrt{2}x_1x_2, x_2^2)$$

$$\phi(\mathbf{z}) = (z_1^2, \sqrt{2}z_1z_2, z_2^2)$$

when learning a linear SVM, recall that the dual solution depends **only on the inner products** of the training data, so we only need to compute inner products in the higher-dimensional space:

$$\begin{aligned}\phi(\mathbf{x})^T \phi(\mathbf{z}) &= x_1^2 z_1^2 + 2x_1 x_2 z_1 z_2 + x_2^2 z_2^2 \\ &= (x_1 z_1 + x_2 z_2)^2 = (\mathbf{x}^T \mathbf{z})^2 \\ &= (\mathbf{x}^T \mathbf{z})^2\end{aligned}$$



the function $\kappa(\mathbf{x}, \mathbf{z}) = (\mathbf{x}^T \mathbf{z})^2$ is an example of a **kernel function**
the kernel function relates the inner-products in the original and transformed spaces; with a kernel, **we can avoid explicit transformation**

Soft margin linear SVM classifier objective

$$\begin{aligned} & \underset{\mathbf{w}, b, \zeta}{\text{minimize}} \quad \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum_{i=1}^m \zeta^{(i)} \quad \longrightarrow \quad \textbf{Primal form} \\ & \text{subject to} \quad y^{(i)} \left(\mathbf{w}^T \mathbf{x}^{(i)} + b \right) \geq 1 - \zeta^{(i)} \quad \text{and} \quad \zeta^{(i)} \geq 0 \quad \text{for } i = 1, 2, \dots, m \end{aligned}$$

Dual form of the linear SVM objective

$$\begin{aligned} & \underset{\alpha}{\text{minimize}} \quad \frac{1}{2} \sum_{i=1}^m \sum_{j=1}^m \alpha^{(i)} \alpha^{(j)} y^{(i)} y^{(j)} \mathbf{x}^{(i)T} \mathbf{x}^{(j)} \quad - \quad \sum_{i=1}^m \alpha^{(i)} \\ & \text{subject to} \quad \alpha^{(i)} \geq 0 \quad \text{for } i = 1, 2, \dots, m \end{aligned}$$

From the dual solution to the primal solution

$$\widehat{\mathbf{w}} = \sum_{i=1}^m \hat{\alpha}^{(i)} y^{(i)} \mathbf{x}^{(i)}$$

$$\hat{b} = \frac{1}{n_s} \sum_{\substack{i=1 \\ \hat{\alpha}^{(i)} > 0}}^m \left(y^{(i)} - \widehat{\mathbf{w}}^T \mathbf{x}^{(i)} \right)$$

Kernel Trick

In Machine Learning, **a kernel is a function** capable of computing the dot product $\phi(\mathbf{a})^T \phi(\mathbf{b})$ based only on the original vectors \mathbf{a} and \mathbf{b} , without having to compute (or even to know about) the transformation ϕ .

$$K(\mathbf{a}, \mathbf{b}) = \phi(\mathbf{a})^T \phi(\mathbf{b})$$

Example of kernelized SVM

Second-degree polynomial mapping

$$\phi(\mathbf{x}) = \phi\left(\begin{pmatrix} x_1 \\ x_2 \end{pmatrix}\right) = \begin{pmatrix} x_1^2 \\ \sqrt{2} x_1 x_2 \\ x_2^2 \end{pmatrix}$$

Kernel trick for a 2nd-degree polynomial mapping

$$\begin{aligned} \phi(\mathbf{a})^T \phi(\mathbf{b}) &= \begin{pmatrix} a_1^2 \\ \sqrt{2} a_1 a_2 \\ a_2^2 \end{pmatrix}^T \begin{pmatrix} b_1^2 \\ \sqrt{2} b_1 b_2 \\ b_2^2 \end{pmatrix} = a_1^2 b_1^2 + 2a_1 b_1 a_2 b_2 + a_2^2 b_2^2 \\ &= (a_1 b_1 + a_2 b_2)^2 = \left(\begin{pmatrix} a_1 \\ a_2 \end{pmatrix}^T \begin{pmatrix} b_1 \\ b_2 \end{pmatrix} \right)^2 = (\mathbf{a}^T \mathbf{b})^2 \end{aligned}$$

Most frequently used popular Kernels

Linear: $K(\mathbf{a}, \mathbf{b}) = \mathbf{a}^T \mathbf{b}$

Polynomial: $K(\mathbf{a}, \mathbf{b}) = \left(\gamma \mathbf{a}^T \mathbf{b} + r \right)^d$

Gaussian RBF: $K(\mathbf{a}, \mathbf{b}) = \exp \left(-\gamma \| \mathbf{a} - \mathbf{b} \|^2 \right)$

Sigmoid: $K(\mathbf{a}, \mathbf{b}) = \tanh \left(\gamma \mathbf{a}^T \mathbf{b} + r \right)$

Mercer's Theorem

According to *Mercer's theorem*, if a function $K(\mathbf{a}, \mathbf{b})$ respects a few mathematical conditions called *Mercer's conditions* (K must be continuous, symmetric in its arguments so $K(\mathbf{a}, \mathbf{b}) = K(\mathbf{b}, \mathbf{a})$, etc.), then there exists a function ϕ that maps \mathbf{a} and \mathbf{b} into another space (possibly with much higher dimensions) such that $K(\mathbf{a}, \mathbf{b}) = \phi(\mathbf{a})^T \phi(\mathbf{b})$. So you can use K as a kernel since you know ϕ exists, even if you don't know what ϕ is. In the case of the Gaussian RBF kernel, it can be shown that ϕ actually maps each training instance to an infinite-dimensional space, so it's a good thing you don't need to actually perform the mapping!

Note that some frequently used kernels (such as the Sigmoid kernel) don't respect all of Mercer's conditions, yet they generally work well in practice.

Why Support Vector Machines (SVM) is good?

Clear Margin Separation: Input of the SVM is a set of datasets and output is an optimal hyperplane. In one dimensional space, the hyperplane is a point. In two-dimensional space, the hyperplane is a line. In three-dimensional space, the hyperplane is a surface. [7]

Effective in high-dimensional Spaces: When the number of dimensions is greater than the number of samples, SVM is automatically regularized. That avoids the overfitting of the high-dimensional data.[7]

Memory Efficient: Instead of using the entire training set, SVM only uses a subset of the training dataset to train the models. The subset is called support vectors.

Effective in non-linear classification: SVM uses the kernel function to map the data into high-dimensional spaces, then separate them linearly. It can solve complex problems with an appropriate kernel function.

Straightforward classification concept: SVM maximizes the margin between the support vectors and the target hyperplane.